

# TREES

- **Overview**

- [Trees](#)
- [Search Trees](#)

- **References**

- Richard F. Gilberg and Behrouz A. Forouzan: Data Structures - A Pseudocode Approach with C. 2nd Edition. Thomson (2005)
- Russ Miller and Laurence Boxer: Algorithms Sequential & Parallel. 2nd Edition. Charles River Media Inc. (2005)
- Stanley B. Lippman, Josée Lajoie, and Barbara E. Moo: C++ Primer. 5th Edition. Addison-Wesley (2013)
- Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, Clifford Stein, “Introduction to Algorithms”, 4th Edition, The MIT Press (2022)
- Scott Meyers: Effective Modern C++. O'Reilly (2014)
- Anthony Williams: C++ Concurrency in Action - Practical Multithreading. Manning Publications Co. (2012)

# TREE BASICS

- A tree  $T$  is a finite, non-empty set of nodes,

$$T = \{r\} \cup T_1 \cup T_2 \cup \dots \cup T_n,$$

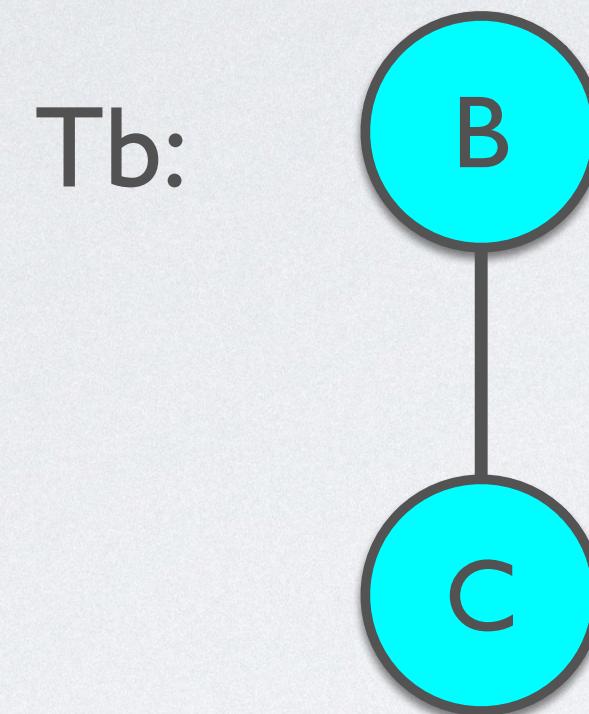
with the following properties:

- A designated node of the set,  $r$ , is called the **root** of the tree.
- The remaining nodes are partitioned into  $n \geq 0$  subsets  $T_1, T_2, \dots, T_n$ , each of which is a **sub-tree**.

# PARENT, CHILD, AND LEAF

- The root node  $r$  of tree  $T$  is the **parent** of all the roots  $r_i$  of the subtrees  $T_i$ ,  $1 < i \leq n$ .
- Each root  $r_i$  of subtree  $T_i$  of tree  $T$  is called a **child** of  $r$ .
- A **leaf** node is a tree with no subtrees.

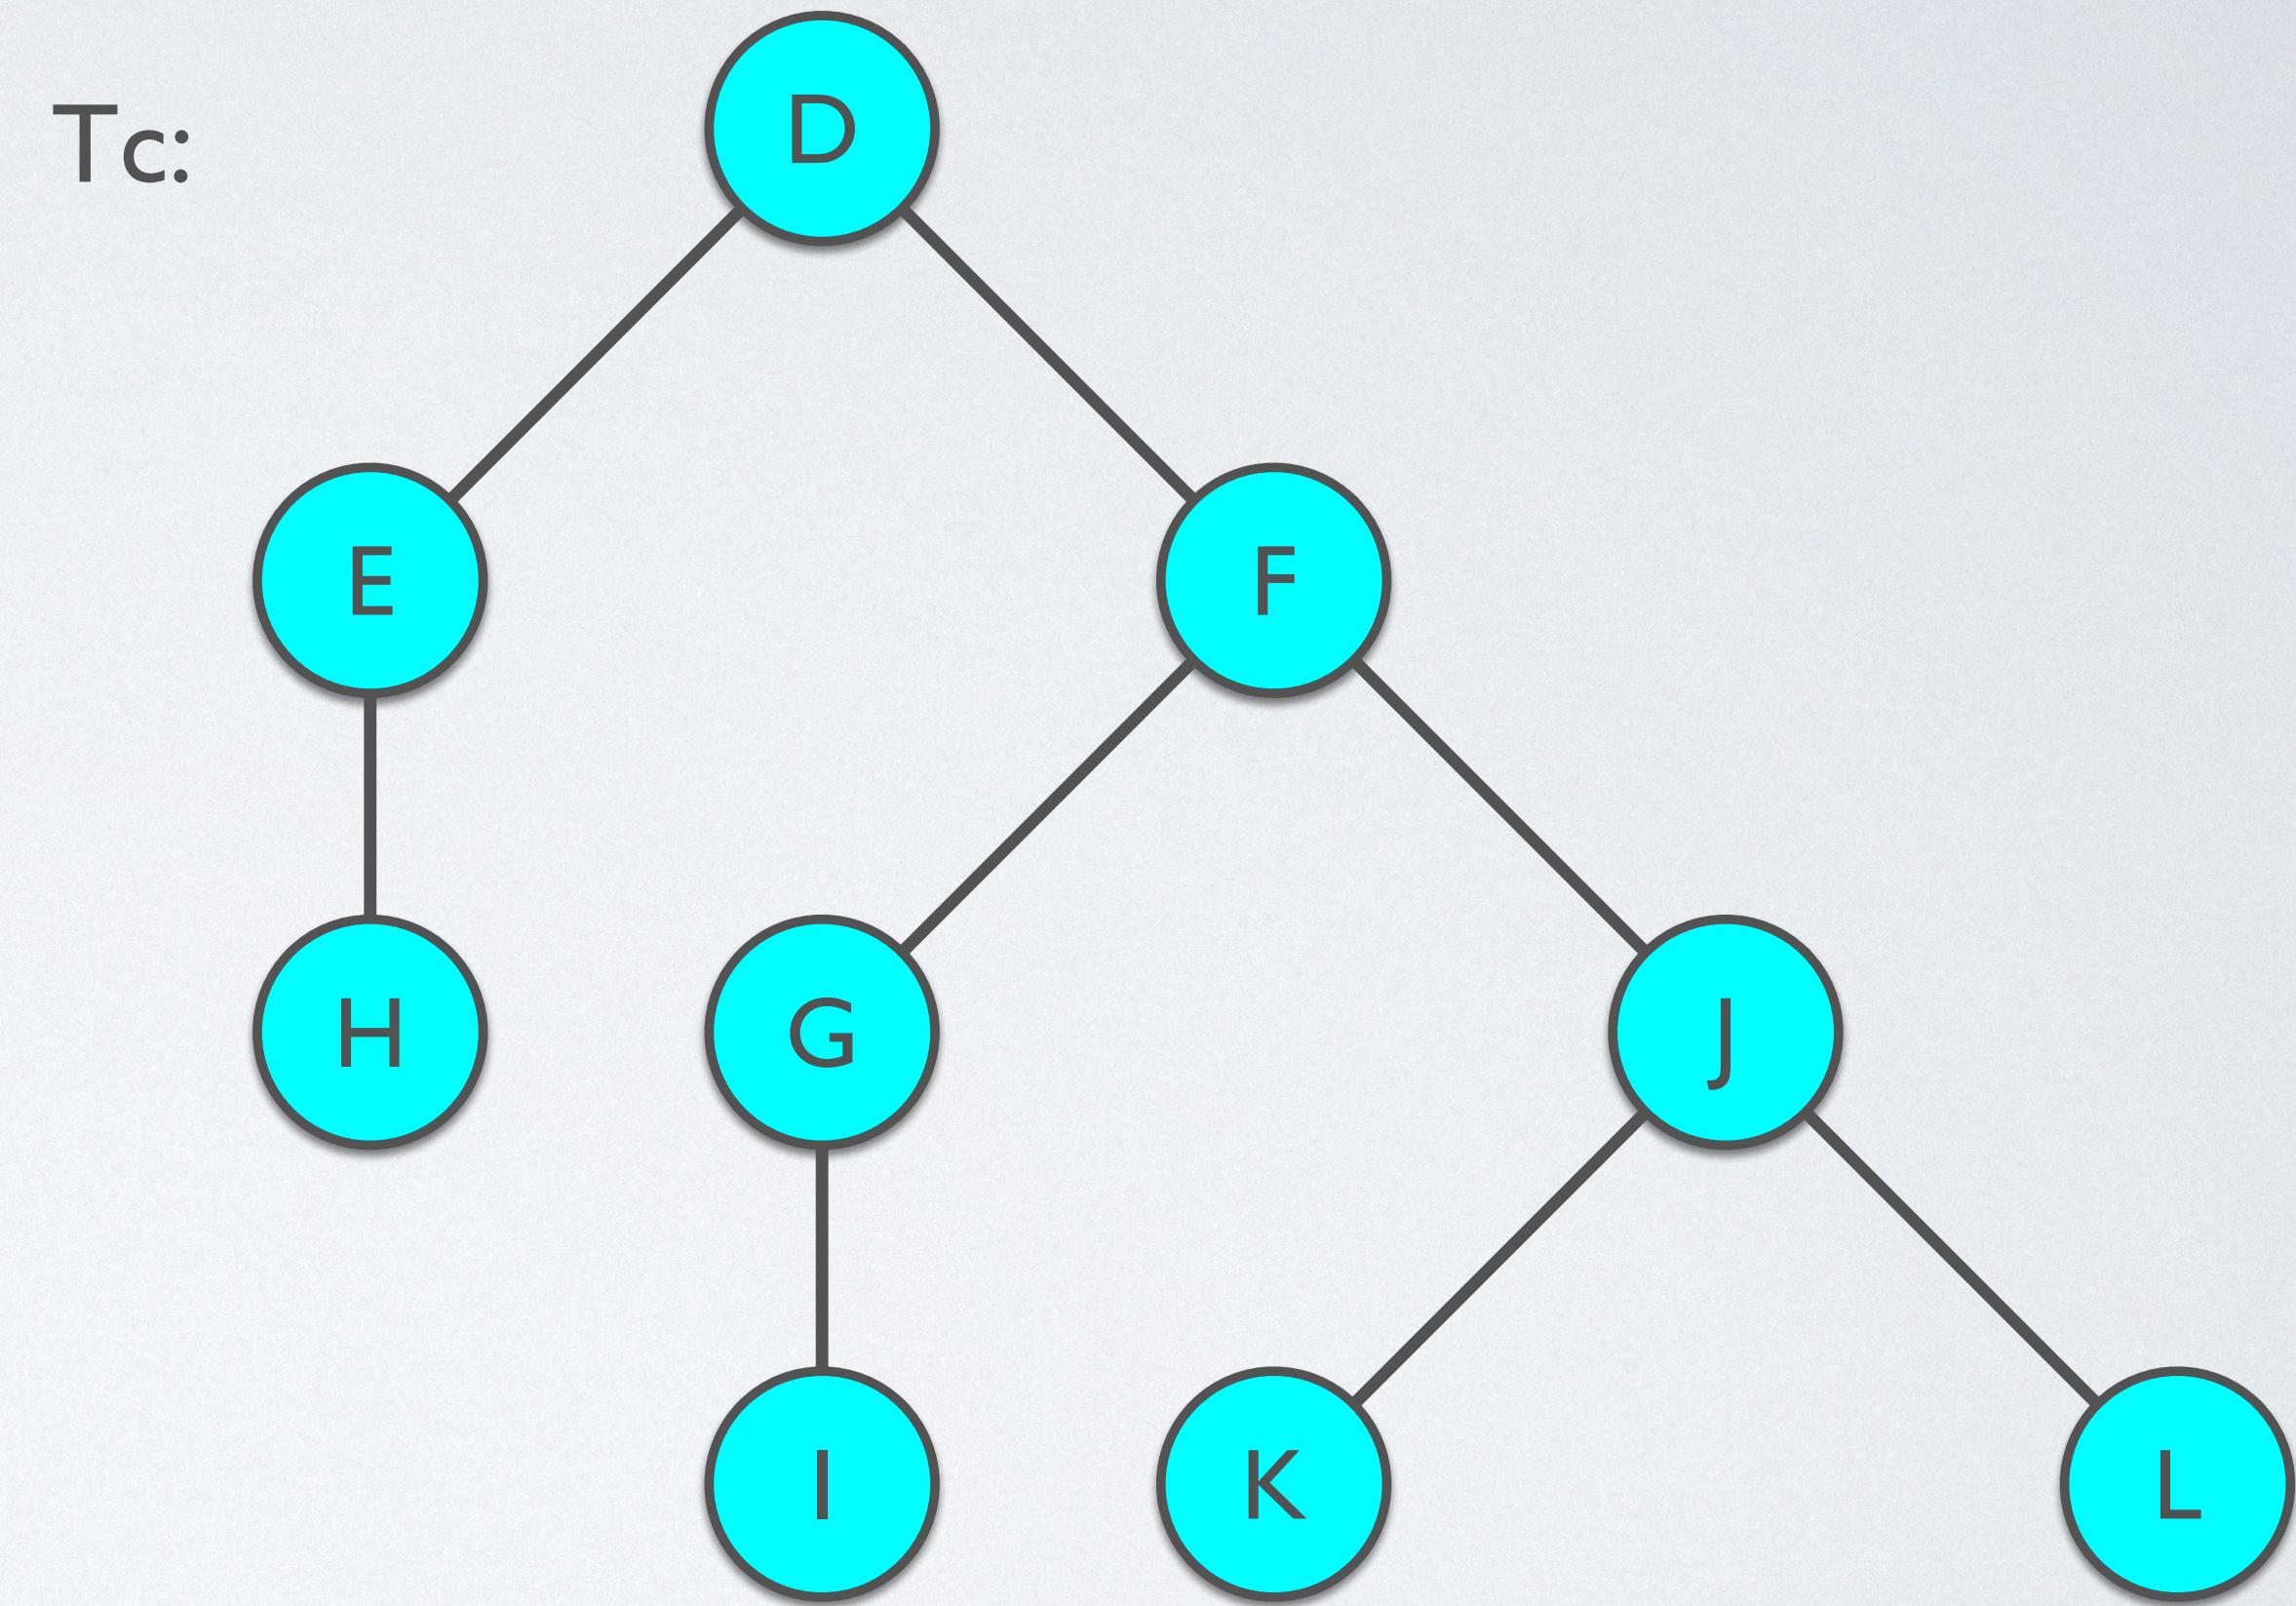
# EXAMPLE TREES



T<sub>a</sub> = {A}

T<sub>b</sub> = {B, {C}}

T<sub>c</sub> = {D, {E, {H}}, {F, {G, {I}}}, {J, {K}, {L}}}}



# DEGREE OF A TREE

- The **degree** of a node is the number of sub-trees associated with that node. For example, the degree of  $T_c = \{D, \{E, \{H\}\}, \{F, \{G, \{I\}\}, \{J, \{K\}, \{L\}\}\}\}$  is 2.
- A node of degree zero has no subtrees. Such a node is called a **leaf**. For example, the leaves of  $T_c$  are H, I, K, L.
- Two roots  $r_i$  and  $r_j$ , of distinct subtrees  $T_i$  and  $T_j$  with the same parent in tree  $T$  are called **siblings**. For example,  $T_i = \{G, \{I\}\}$  and  $T_j = \{J, \{K\}, \{L\}\}$  are siblings in  $T_c$ .

# PATH AND PATH LENGTH

- Given a tree  $T$  containing the set of nodes  $R$ , a **path  $p$**  in  $T$  is defined as a non-empty sequence of nodes

$$p = \{r_1, r_2, \dots, r_k\}$$

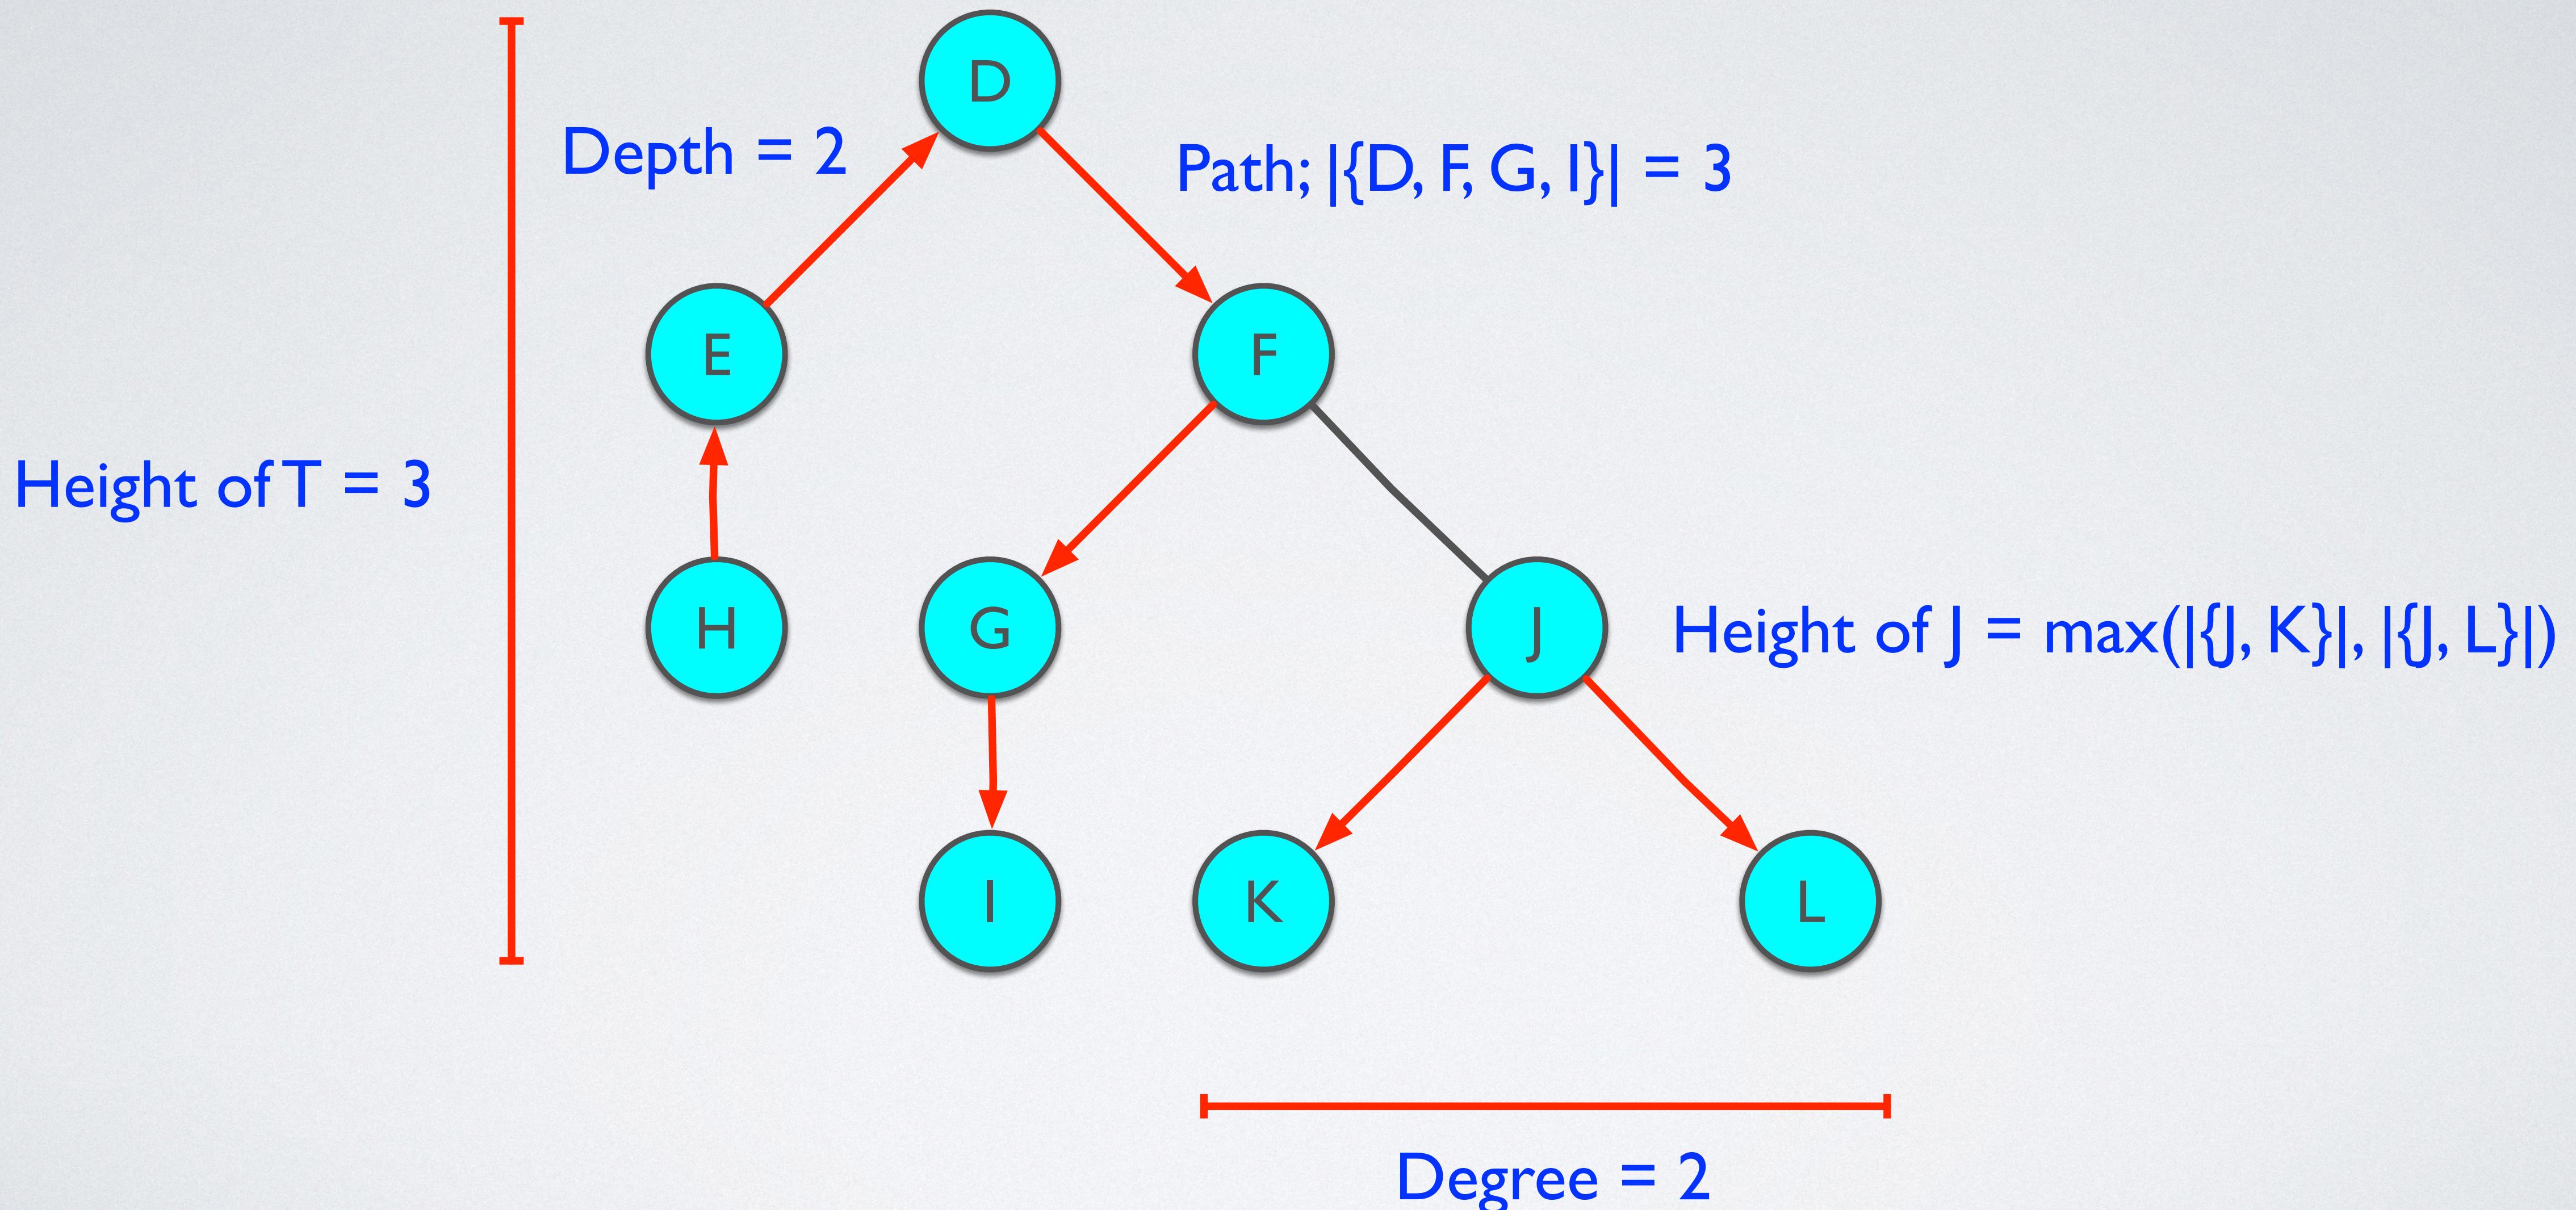
where  $r_i \in R$ , for  $1 \leq i \leq k$  such that the  $i$ th node in the sequence,  $r_i$ , is the parent of the  $(i+1)$ th node in the sequence,  $r_{i+1}$ .

- The **length of path  $p$** ,  $|p|$ , is  $k-1$ , which corresponds to the distance from the root  $r_1$  to the leaf  $r_k$  in tree  $T$ .

# DEPTH AND HEIGHT

- The **depth** of a node  $r_i \in R$  in a tree  $T$  is the length of the unique path  $p$  in  $T$  from its root to the node  $r_i$ .
- The **height of a node**  $r_i \in R$  in a tree  $T$  is the length of the longest path from node  $r_i$  to a leaf. Consequently, all leaves have a height zero.
- The **height of a tree**  $T$  is the **height of its root node**  $r$ .

# DEGREE, PATH, DEPTH, AND HEIGHT



# TREE NODES WITH THE SAME DEGREE

- The general case allows each node in a tree to have a different degree.
- We now consider a variant of trees in which each node has the same degree.
- Unfortunately, it is impossible to construct a tree with a finite number of nodes that all have the same degree  $N$ , except the trivial case  $N = 0$ .
- We need a special notation, called empty tree, to realize trees in which all nodes have the same degree.

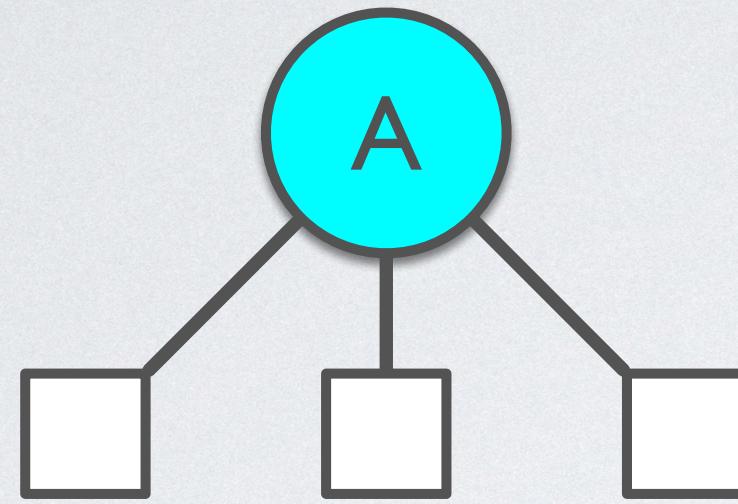
# N-ARY TREES

- An N-ary tree  $T$ ,  $N \geq 1$ , is a finite set of nodes with one of the following properties:
  - Either the set is empty,  $T = \emptyset$ , or
  - The set consists of a root,  $R$ , and exactly  $N$  distinct  $N$ -ary trees. That is, the remaining nodes are partitioned into  $N \geq 1$  subsets,  $T_1, T_2, \dots, T_N$ , each of which is an  $N$ -ary tree such that

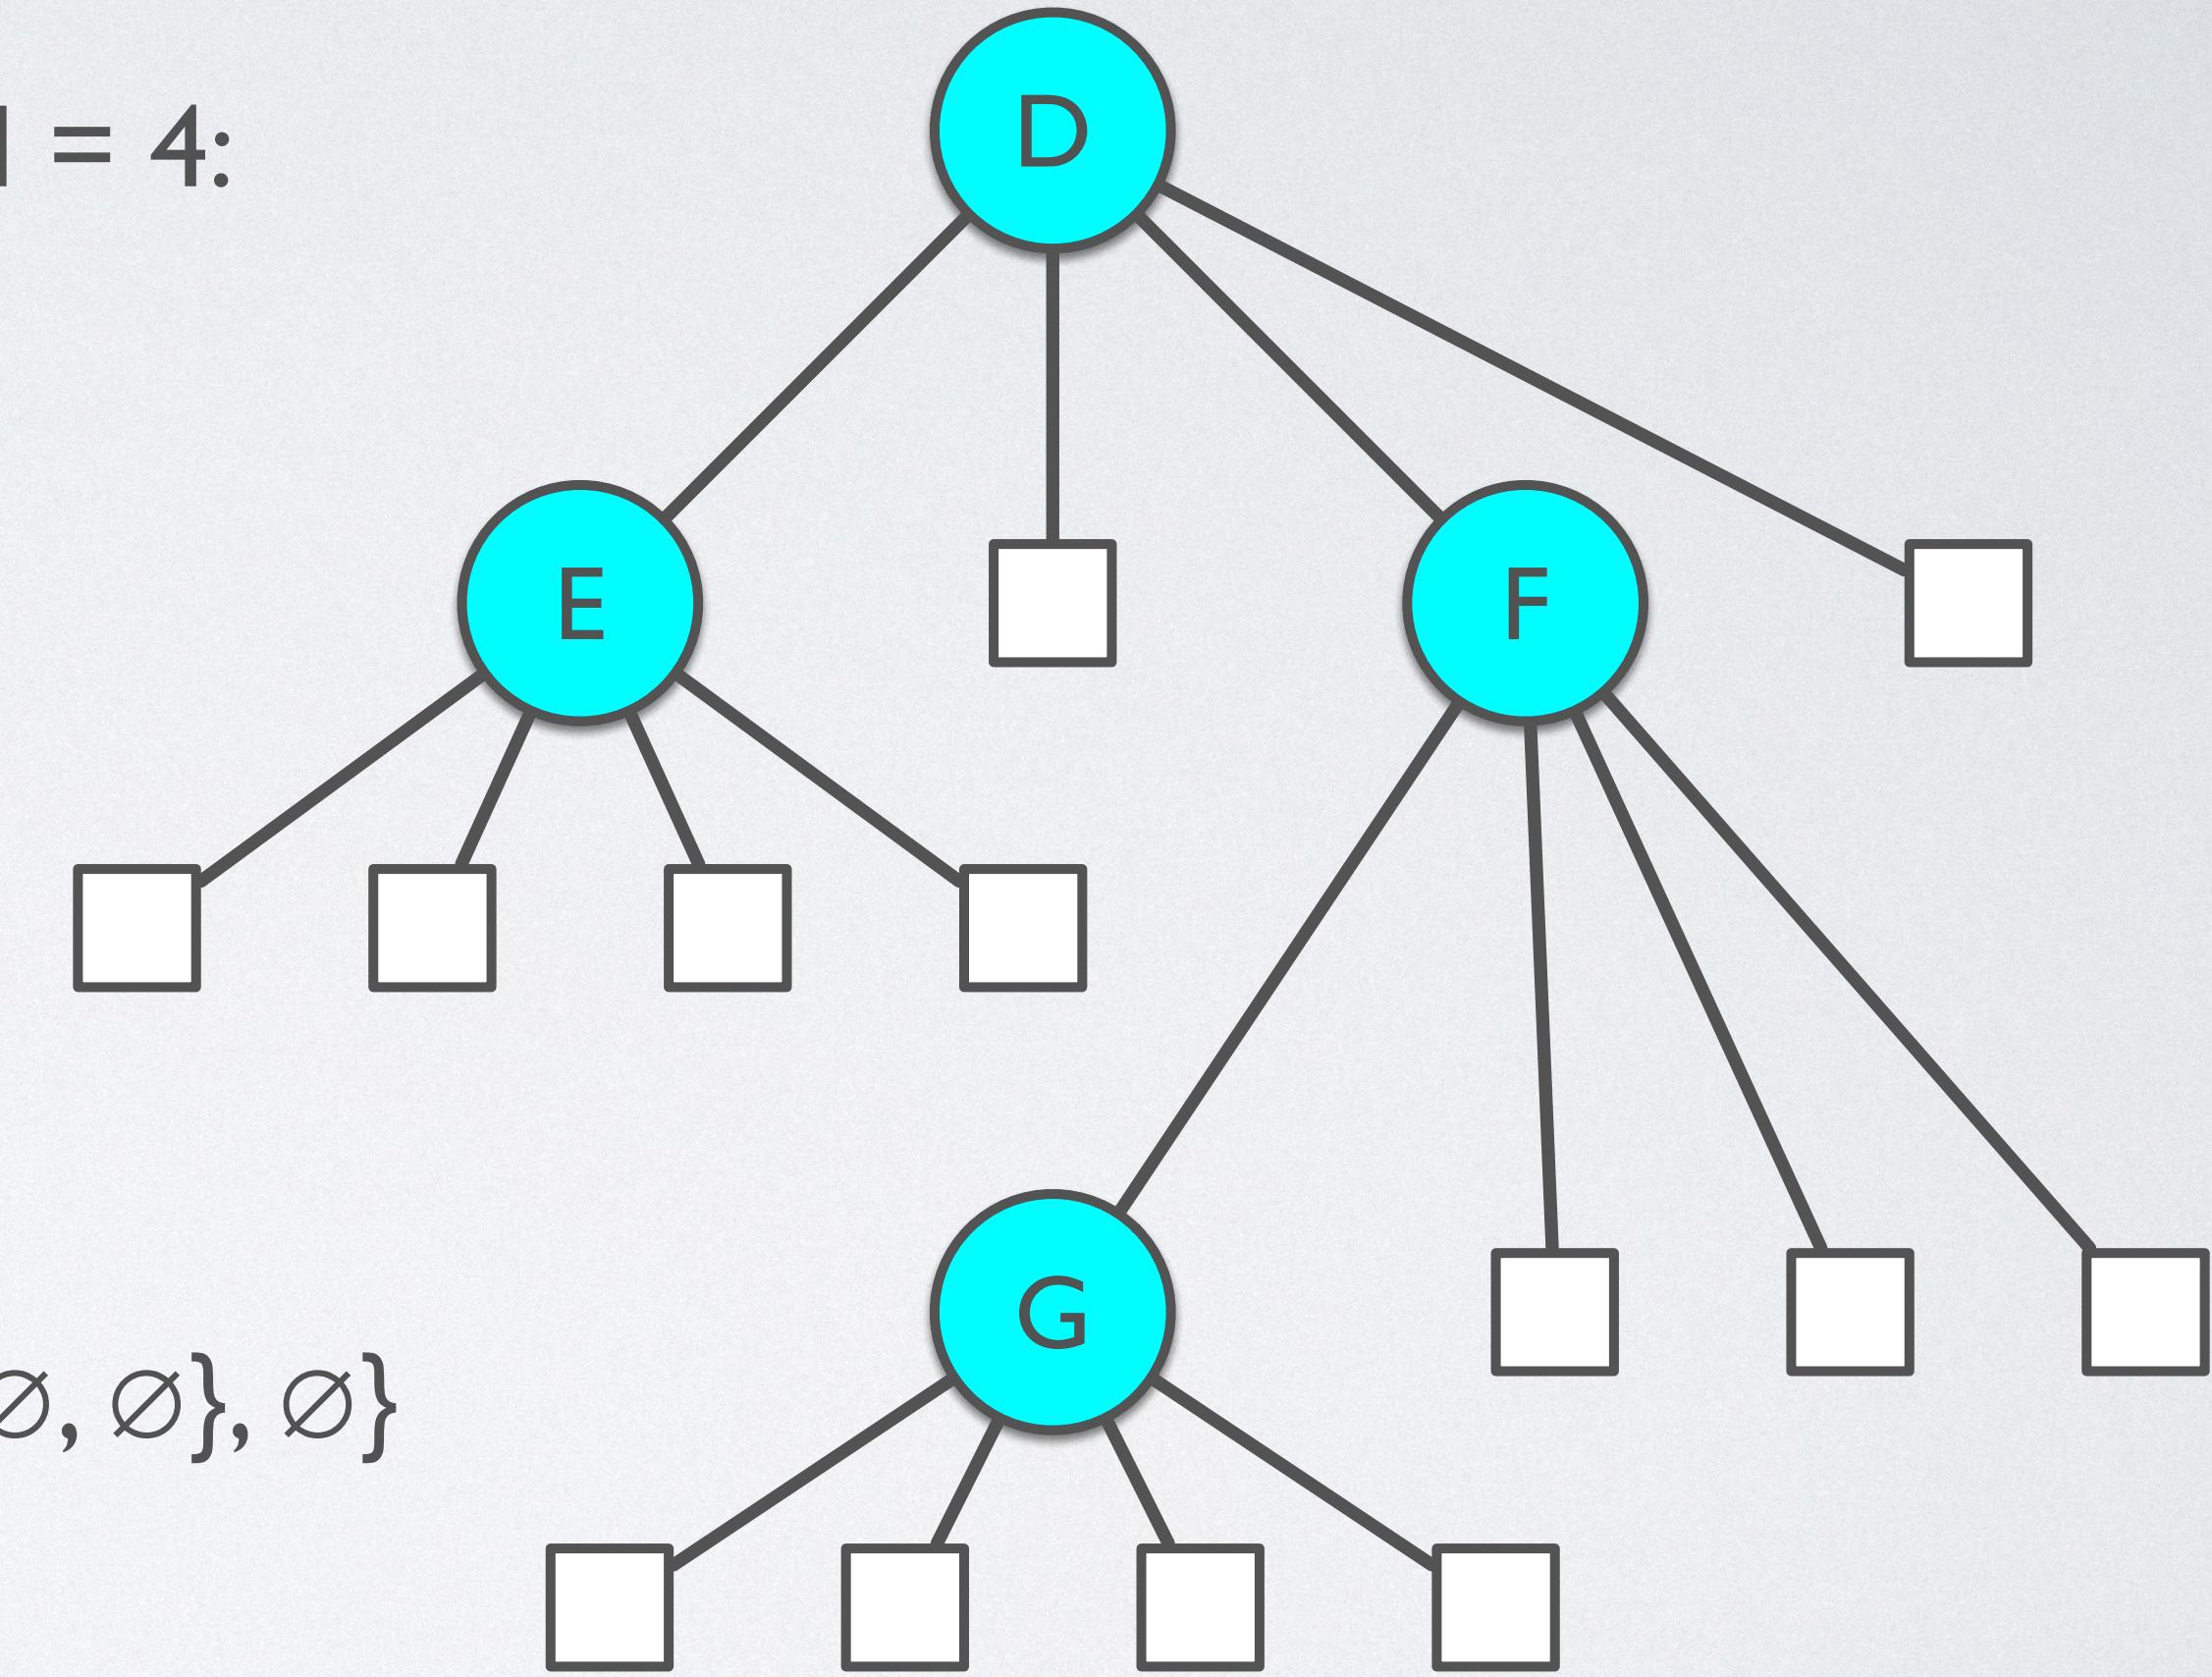
$$T = \{R, T_1, T_2, \dots, T_N\}.$$

# EXAMPLE N-ARY TREES

T<sub>a</sub>; N = 3:



T<sub>b</sub>; N = 4:



T<sub>a</sub> = {A,  $\emptyset$ ,  $\emptyset$ ,  $\emptyset$ }

T<sub>b</sub> = {D, {E,  $\emptyset$ ,  $\emptyset$ ,  $\emptyset$ ,  $\emptyset$ },  $\emptyset$ , {F, {G,  $\emptyset$ ,  $\emptyset$ ,  $\emptyset$ ,  $\emptyset$ },  $\emptyset$ ,  $\emptyset$ ,  $\emptyset$ },  $\emptyset$ }

# THE EMPTY TREE

- The empty tree,  $T = \emptyset$ , is a tree.
- From the modeling point of view, an empty N-ary tree:
  - has no key, and
  - has to have the same type as a non-empty N-ary tree.
- Using null (i.e., **nullptr**) to denote an empty N-ary tree may be unsuitable, as null refers to nothing!

# SENTINEL NODE

- A sentinel node is a programming idiom used to facilitate tree-based operations.
- A sentinel node in tree structures indicates a node with no children.
- Sentinel nodes behave like null-pointers. However, unlike null-pointers, which refer to nothing, sentinel nodes denote proper, yet empty, subtrees.
- Smart pointers allow us to represent sentinel nodes and non-empty trees equally. In particular, for an N-ary tree, we employ unique smart pointers (i.e., `std::unique_ptr`) to represent N-ary tree nodes. In the case of the sentinel node, the smart pointer wraps **nullptr**, whereas for non-empty N-ary trees, it wraps a unique, ownership-managed heap location.

```

template<typename T, size_t N>
class NTree
{
public:
    using node = std::unique_ptr<NTree>;

    NTree( const T& aKey = T{} ) noexcept;
    NTree( T&& aKey ) noexcept;

    template<typename... Args>
    static Node makeNode( Args&&... args );

    NTree( const NTree& aOther );
    NTree& operator=( const NTree& aOther );

    NTree( NTree&& aOther ) noexcept;
    NTree& operator=( NTree&& aOther ) noexcept;

    void swap( NTree& aOther ) noexcept;
    const T& operator*() const noexcept;
    NTree& operator[]( size_t alndex ) const;

    void attach( size_t alndex, node& aNode );
    node detach( size_t alndex );

    bool leaf() const noexcept;
    size_t height() const noexcept;

private:
    T fKey;
    node fNodes[N];
};

```

# NTREE<T, N>

- Class NTree is parameterized over the key type **T** and the degree of the tree **N**.
- Objects of class **NTree** do not support shared ownership. Nodes take exclusive ownership.
- Attaching and detaching sub-trees transfers ownership.
- The indexer returns a reference to the corresponding node. The ownership of a node remains with the tree, but we can obtain the address of the returned node.
- **NTree** defines proper copy and move semantics. The synthesized destructor works as expected.
- The auxiliary methods leaf() and height() test a node for the leaf property and return the height, respectively.

# NTREE<T,N> CONSTRUCTORS

- We can construct an N-ary tree node by copying the key value or by moving the key value.

```
NTree( const T& aKey = T{} ) noexcept :  
    fKey(aKey)
```

```
{}
```

```
NTree( T&& aKey ) noexcept :  
    fKey(std::move(aKey))
```

```
{}
```

- In the lvalue constructor, we can initialize an N-ary tree node with a default key, T{}. Hence, this constructor can be used as the default constructor.
- The constructors create leaf nodes. All sub-trees are initialized as sentinel nodes through the default initialization of std::unique\_ptr.

# MAKENODE METHOD

```
template<typename... Args>
static Node makeNode( Args&&... args )
{
    return std::make_unique<NTree>( std::forward<Args>(args)... );
}
```

- The static method `makeNode()` creates N-ary tree objects with unique ownership semantics.
- We use a variadic template method and perfect forwarding to allow for different ways to create N-ary tree nodes. In particular, `makeNode()` can create new N-ary tree nodes from keys and other N-ary tree nodes using copy and move semantics.

# KEY AND SUB-TREE ACCESS

```
const T& operator*() const noexcept
```

```
{  
    return fKey;  
}
```

```
NTree& operator[]( size_t alndex ) const
```

```
{  
    // valid index and node  
    assert( alndex < N && fNodes[alndex] );  
  
    return *fNodes[alndex];  
}
```

- The dereference operator returns the key value as a constant reference.
- The node indexer returns a reference to the indexed sub-tree. Ownership of this sub-tree remains unaffected.
- The indexer checks that we select a valid tree node (i.e., within range of N and not a sentinel node).
- **Note:** It is impossible to obtain a reference to **nullptr**. Hence, the indexer cannot return a reference to a sentinel node. But this has the desired effect.

# ATTACHING A SUB-TREE

```
void attach( size_t alndex, Node& aNode )
{
    // transfers ownership from aNode to fNodes[alndex]

    assert( alndex < N&& !fNodes[alndex] );

    fNodes[alndex] = std::move(aNode);
}
```

- To attach an N-ary tree node to a given target node, we require a valid sub-tree index (i.e., an index occupied by a sentinel node) and transfer ownership of the sub-tree to the target node.
- Unique smart pointers can only be moved, not copied. Ownership of the managed heap space is transferred (i.e., after method attach() finishes, aNode is an empty N-ary tree node).

# DETACHING A SUB-TREE

```
Node detach( size_t alndex )
{
    // transfers ownership from fNodes[alndex] to result

    assert( alndex < N && fNodes[alndex] );

    return std::move(fNodes[alndex]);
}
```

- To detach an N-ary tree node from a given source node, we require a valid sub-tree index (i.e., an index occupied by an N-ary tree node) and transfer ownership of the sub-tree to the caller.
- After method detach() finishes, the N-ary sub-tree at index alndex is an empty N-ary tree node (i.e., a sentinel node).

```

bool leaf() const noexcept
{
    for ( size_t i = 0; i < N; i++ )
    {
        if ( fNodes[i] )
        {
            return false;
        }
    }

    return true;
}

size_t height() const noexcept
{
    size_t Result = 0;

    if ( !leaf() )
    {
        for ( size_t i = 0; i < N; i++ )
        {
            if ( fNodes[i] )
            {
                Result = std::max( Result, fNodes[i]->height() + 1 );
            }
        }
    }

    return Result;
}

```

# LEAF() AND HEIGHT()

- The method `leaf()` returns true if all sub-trees of the current node are sentinels. The running time is  $O(n)$ .
- The method `height()` returns the height of the current node, which is the longest path from the current node to a leaf. Leaf nodes have height zero. The running time is  $O(n)$ .

```

NTree( const NTree& aOther ) :
    fKey(aOther.fKey)
{
    for ( size_t i = 0; i < N; i++ )
    {
        if ( aOther.fNodes[i] )
        {
            // copy non-empty subtree
            fNodes[i] = std::move(makeNode( *aOther.fNodes[i] ));
        }
    }
}

NTree& operator=( const NTree& aOther )
{
    if ( this != &aOther )
    {
        this->~NTree();
        new (this) NTree( aOther );
    }

    return *this;
}

```

# COPY SEMANTICS

- The copy constructor for N-ary trees first copies the root key and then recursively copies all non-empty sub-trees. The ownership of the copied sub-trees is transferred to the newly created N-ary tree node.
- The copy assignment operator is defined in the standard way, first releasing all resources and then performing an in-place copy construction.

# MOVE SEMANTICS

- The move operations are defined in the standard way and rely on the existence of a suitable swap() method.
- Once a move operation has been completed, the source node, aOther, is an empty tree.

```
NTree( NTree&& aOther ) noexcept :  
    NTree()  
{  
    swap( aOther );  
}  
  
NTree& operator=( NTree&& aOther ) noexcept  
{  
    if ( this != &aOther )  
    {  
        swap( aOther );  
    }  
  
    return *this;  
}
```

# SWAP FOR N-ARY TREES

```
void swap( NTree& aOther ) noexcept
{
    std::swap( fKey, aOther.fKey );

    for ( size_t i = 0; i < N; i++ )
    {
        std::swap( fNodes[i], aOther.fNodes[i] );
    }
}
```

- The method `swap` exchanges the root keys and the sub-trees.
- The ownership of the sub-trees is mutually transferred.
- The method `std::swap()` internally uses move semantics to and from a temporary variable. This guarantees that unique smart pointers can be swapped, that is, their heap locations and the ownership can be mutually exchanged.

```

std::string a( "root" );
std::string a1( "a1" );
std::string a2( "a2" );
std::string a3( "a3" );
std::string a12( "a12" );

using NS3Tree = NTree<std::string, 3u>;
using NS3TreeNode = NS3Tree::Node;

NS3TreeNode root = NS3Tree::makeNode( a );
NS3TreeNode n1 = NS3Tree::makeNode( a1 );
NS3TreeNode n2 = NS3Tree::makeNode( a2 );
NS3TreeNode n3 = NS3Tree::makeNode( a3 );
NS3TreeNode n12 = NS3Tree::makeNode( a12 );

root->attach( 0u, n1 );
root->attach( 1u, n2 );
root->attach( 2u, n3 );
(*root)[1].attach( 1u, n12 );

std::cout << "Constructed root:" << std::endl;

std::cout << "root:\t\t" << **root << std::endl;
std::cout << "root[0]:\t" << *(*root)[0u] << std::endl;
std::cout << "root[1]:\t" << *(*root)[1u] << std::endl;
std::cout << "root[2]:\t" << *(*root)[2u] << std::endl;
std::cout << "root[1][1]:\t" << *(*root)[1u][1u] << std::endl;

std::cout << "height root:\t\t" << root->height() << std::endl;
std::cout << "height root[0]:\t\t" << (*root)[0u].height() << std::endl;
std::cout << "height root[1]:\t\t" << (*root)[1u].height() << std::endl;
std::cout << "height root[2]:\t\t" << (*root)[2u].height() << std::endl;
std::cout << "height root[1][1]:\t\t" << (*root)[1u][1u].height() << std::endl;

std::cout << "Move root:" << std::endl;

NS3TreeNode move = NS3Tree::makeNode( std::move(*root) );

std::cout << "height root:\t\t" << root->height() << std::endl;

```

# NTREE<T,N> EXAMPLE

Microsoft Visual Studio Deb...

```

Constructed root:
root:          root
root[0]:       a1
root[1]:       a2
root[2]:       a3
root[1][1]:   a12
height root:    2
height root[0]:  0
height root[1]:  1
height root[2]:  0
height root[1][1]: 0
Move root:
height root:    0

```

# 2-ARY TREES: BINARY TREES

- A binary tree  $T$  is a finite set of nodes with one of the following properties:
  - Either the set is empty,  $T = \emptyset$ , or
  - The set consists of a root,  $r$ , and exactly 2 distinct binary trees  $T_L$  and  $T_R$ ,  $T = \{r, T_L, T_R\}$ .
  - The tree  $T_L$  is called the **left subtree** of  $T$  and the tree  $T_R$  is called the **right subtree** of  $T$ .

# SEMANTIC MISMATCH

```
template<typename T>
using BTree<T> = NTree<T,2>;
```

- Representing a binary tree simply as a type alias to `NTree<T,2>` works.  
Unfortunately, this type alias does not convey the intended semantics of a  
binary tree, namely that a binary tree has a dedicated left and right sub-tree.

```

template<typename T>
class BTree
{
public:
    using node = std::unique_ptr<BTree>;

    BTree( const T& aKey = T{} ) noexcept;
    BTree( T&& aKey ) noexcept;

    template<typename... Args>
    static node makeNode( Args&&... args );

    BTree( const BTree& aOther );
    BTree& operator=( const BTree& aOther );

    BTree( BTree&& aOther ) noexcept;
    BTree& operator=( BTree&& aOther ) noexcept;

    void swap( BTree& aOther ) noexcept;

    const T& operator*() const noexcept;
    bool hasLeft() const noexcept;
    BTree& left() const;
    bool hasRight() const noexcept;
    BTree& right() const;

    void attachLeft( node& aNode );
    void attachRight( node& aNode );

    node detachLeft();
    node detachRight();

    bool leaf() const noexcept;
    size_t height() const noexcept;

private:
    T fKey;
    node fLeft;
    node fRight;
};

}

```

# BTREE<T>

- The template class BTree<T> defines the semantics of a binary tree.
- BTree<T> provides dedicated access to the left and right sub-tree of a given binary tree node.
- The implementation follows that for N-ary tree with the exception that BTree<T> defines two predicates: `hasLeft()` and `hasRight()`, which return **true** if the corresponding sub-tree is not an empty tree. **These predicates are required for tree traversal.**
- We cannot simply use subclassing and define BTree<T> as a subtype of NTree<T, 2>. The methods `hasLeft()` and `hasRight()` require access to the sub-tree nodes, which are private in NTree. Exposing the sub-trees to sub-classes is not safe.